





# Approximation Algorithms

---




Hardest

Hard

Medium

Easy

# NP-Completeness



"I can't find an efficient algorithm, but neither can all these famous people."



# ***Coping With NP-Hardness***

---

## **Brute-force Algorithms.**

- Develop clever enumeration strategies.
- Guaranteed to find optimal solution.
- No guarantees on the running time.

## **Heuristics.**

- Develop intuitive algorithms.
- Guaranteed to run in polynomial time.
- No guarantees on the quality of solution.

## **Approximation Algorithms.**

- Guaranteed to run in polynomial time.
- Guaranteed to find "high quality" solution, within 1% of optimum.

*Obstacle: We need to prove a solution's value is close to the optimum, without even knowing what optimum value is!*

## Coping with NP-completeness

---

**Q.** Suppose I need to solve an **NP**-hard optimization problem. What should I do?

**A.** Sacrifice one of three desired features.

- i. Runs in polynomial time.
- ii. Solves arbitrary instances of the problem.
- iii. Finds optimal solution to problem.

$\rho$ -approximation algorithm.

- Runs in polynomial time.
- Solves arbitrary instances of the problem
- Finds solution that is within ratio  $\rho$  of optimum.

**Challenge.** Need to prove a solution's value is close to optimum, without even knowing what is optimum value.



# Approximation Algorithms

---

- The best algorithm for solving the NP-complete problem requires the exponential time in the worst-case. It is too time-consuming.
- To reduce the time required for solving a problem, we can relax the problem, and obtain a feasible solution “close” to optimal solution.



# Approximation Algorithms

---

- One compromise is to use the “Heuristic” Solutions.
- “Heuristic” may be interpreted as an “Educated Guess”.
- Approximation Algorithms return Near Optimal Solutions.
- Need to find an Approximation Ratio Bound for algorithm.



# Approximation Ratio Bound

We say an approximation algorithm for the problem has a ratio bound of  $\rho(n)$  if for any input size  $n$ , the cost  $C$  of the solution produced by the approximation algorithm is within a factor of  $\rho(n)$  of the  $C^*$  of the optimal solution:

$$\max\left\{\frac{C}{C^*}, \frac{C^*}{C}\right\} = \rho(n)$$

This applies for both minimization and maximization problems.





# ***Performance Guarantees***

---

- An approximation algorithm is bounded by  $\rho(n)$  if, for all input of size  $n$ , the cost  $c$  of the solution obtained by an approximation algorithm is within a factor  $\rho(n)$  of  $c^*$  of an optimal solution.



# $\rho$ -approximation algorithm

---

- An approximation algorithm with an approximation ratio bound of  $\rho$  is referred to as  $\rho$ -approximation algorithm or also known as  $(1+\varepsilon)$ -approximation algorithm.
- Note that  $\rho$  is always  $> 1$  and  $\varepsilon = \rho - 1$ .






# Vertex Cover

**Vertex Cover:** a subset of vertices which “**covers**” every edge.  
An edge is **covered** if one of its endpoint is chosen.

## The Minimum Vertex Cover Problem:

Find a vertex cover with minimum number of vertices.









# Vertex Cover Problem

---

- Let  $G=(V, E)$ . The subset  $S$  of  $V$  that meets every edge of  $E$  is known as the **Vertex Cover**.
- The Vertex Cover Problem is solved for finding a vertex cover of the **Minimum** size. It is NP-Hard Problem or Optimization Problem version of an NP-Complete Decision Problem.

# Examples of Vertex Cover





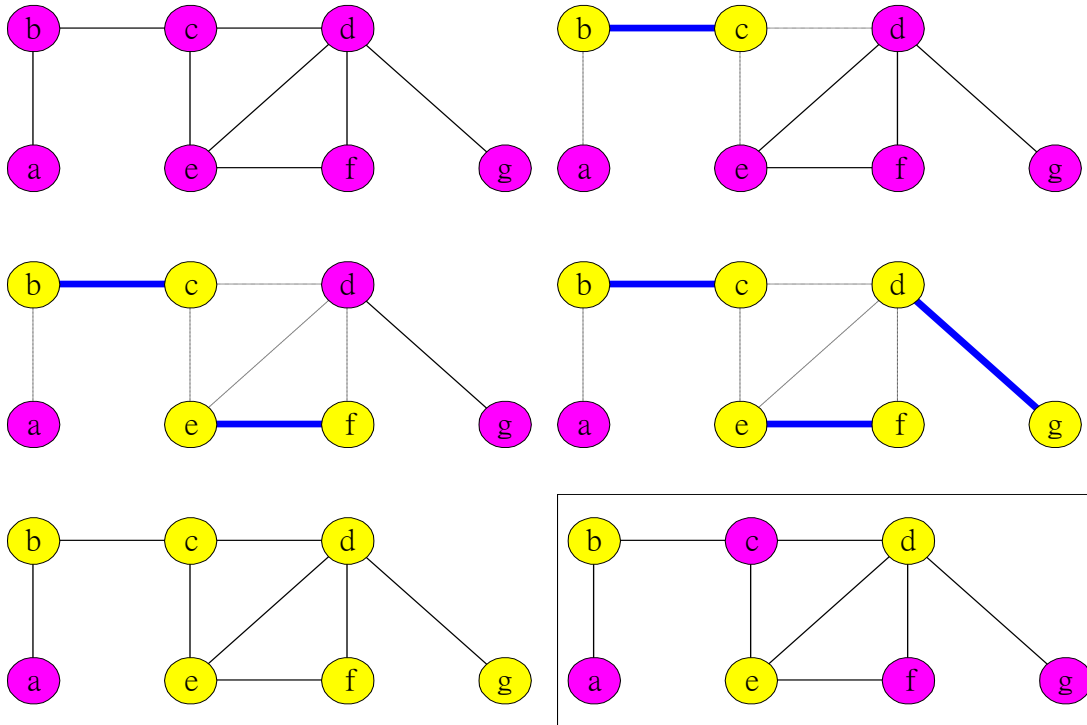
# Vertex Cover Problem

---

APPROX\_VERTEX\_COVER( $G$ )

```
1   $C \leftarrow \phi$ 
2   $E' \leftarrow E(G)$ 
3  while  $E' \neq \phi$ 
4      do let  $(u, v)$  be an arbitrary edge of  $E'$ 
5           $C \leftarrow C \cup \{u, v\}$ 
6          remove from  $E'$  every edge incident on either  $u$  or  $v$ 
7  return  $C$ 
```

# Vertex Cover Problem



Complexity:  $O(E)$



# Vertex Cover Problem

**Theorem:** APPROX\_VERTEX\_COVER has ratio bound of 2.

**Proof.**

$C^*$ : optimal solution

$C$ : approximate solution

$A$ : the set of edges selected in step 4

Let  $A$  be the set of selected edges.

$|C| = 2|A|$  When one edge is selected, 2 vertices are added into  $C$ .

$|A| \leq |C^*|$  No two edges in  $A$  share a common endpoint.

$$\Rightarrow |C| \leq 2|C^*|$$





# Traveling Salesperson Problem

*Traveling Salesperson Problem* (TSP) asks for the shortest Hamiltonian cycle in a weighted undirected graph.

Consider  $G$  be an arbitrary undirected graph with  $n$  vertices.

Length Function  $l(e) = \begin{cases} 1 & \text{if } e \text{ is an edge in } G \\ 2 & \text{otherwise} \end{cases}$  for  $K_n$

Where,  $G$  has a Hamiltonian cycle then there is an Hamiltonian cycle in  $K_n$  whose length is exactly  $n$

*Traveling Salesperson Problem* is NP-hard (NP-complete) even if all the edge lengths are 1 or 2 due to polynomial time reduction from Hamiltonian cycle to this type of Traveling salesperson problem.



# Traveling Salesperson Problem

We can replace the values in length function by any values we like

Length Function 
$$l(e) = \begin{cases} 1 & \text{if } e \text{ is an edge in } G \\ n & \text{otherwise} \end{cases}$$

$G$  has a Hamiltonian cycle then there is an Hamiltonian cycle in  $K_n$  whose length is exactly  $n$  or has length at least  $2n$ .

*Thus if we can approximate* the shortest traveling salesman tour within a factor of 2 in polynomial time we would have a polynomial time algorithm for the Hamiltonian cycle problem

For any function  $f(n)$  that can be computed in polynomial in  $n$ , there is no polynomial time  $f(n)$  approx to TSP on general weighted graph unless  $P=NP$ .



# *TSP: A Special Case*

Edge lengths satisfy triangular inequality  
$$l(u,v) \leq l(u,w) + l(w,v)$$

This is true for geometric graph


- Compute Minimum Spanning Tree  $T$  of the weighted input graph
- Depth First Traversal (Depth First Search) of the MST  $T$
- Numbering the vertices in order that we first encounter them
- Return the cycle found by visiting vertices as per this numbering



# *TSP: A Special Case*

---


Demonstration



Set of points distributed in 2D

# *TSP: A Special Case*


Demonstration



Minimum Spanning Tree


# *TSP: A Special Case*

Demonstration



# *TSP: A Special Case*


Demonstration



Depth First Traversal and Numbering of Vertices

# *TSP: A Special Case*

Demonstration




Traveling Salesperson Tour



# *TSP: A Special Case*


Demonstration



Traveling Salesperson Tour with Cost 2.MST

# *TSP: A Special Case*

Demonstration



Traveling Salesperson Tour with Reduced Cost  $\leq 2 \cdot \text{MST}$



# *TSP: A Special Case*

---

## **Output Quality :**

Cost of the tour using this algorithm


$\leq 2^*$  cost of minimum spanning tree

$\leq 2^*$  cost of optimal solution

Conclusion: The algorithm outputs 2 approximation of the minimum traveling salesperson problem

# *TSP: An Improved Heuristic*


Locate odd degree vertices in minimum spanning tree (MST)



Number of odd degree vertices is even

# *TSP: An Improved Heuristic*


Locate odd degree vertices in minimum spanning tree (MST)



Perfect matching of odd degree vertices

# *TSP: An Improved Heuristic*

Locate odd degree vertices in minimum spanning tree (MST)



Merging the perfect edges with MST




# Bin Packing Problem

---

- Given  $n$  items of sizes  $a_1, a_2, \dots, a_n$ ,  $0 < a_i \leq 1$  for  $1 \leq i \leq n$ , which are to be placed in bins of unit capability, then the bin packing problem can be solved by determining the minimum number of bins to accommodate all the items.
- Consider the items of different sizes with lengths of time of executing different jobs on a standard processor, then we need to use minimum number of processors which can finish all the jobs within a fixed time. // Assume that the longest job takes one unit time i.e. equal to fixed time.

# Example of Bin Packing Problem

- Ex. Given  $n = 5$  items with sizes 0.3, 0.5, 0.8, 0.2, 0.4, then the Optimal Solution is 3 Bins.



The Bin Packing Problem is NP-Hard Optimization Problem.





# An Approximation Algorithm for the Bin Packing Problem

---

- An Approximation Algorithm: (First-Fit (FF)) place the item  $i$  into the lowest-indexed bin which can accommodate the item  $i$ .
- OPT: The number of bins of the Optimal Solution
- FF: The number of bins in the First-Fit Algorithm
- $C(B_i)$ : The sum of the sizes of items packed in the bin  $B_i$  in the First-Fit Algorithm
- Let  $FF=m$ .

# An Approximation Algorithm for the Bin Packing Problem

- $OPT \geq \left\lceil \sum_{i=1}^n a_i \right\rceil$ , ceiling of sum of sizes of all items
- ~~$C(B_i) + C(B_{i+1}) > 1$  (a)(Otherwise, the items in  $B_{i+1}$  will be put in  $B_i$ ).~~  $C(B_i)$ : The sum of sizes of items packed in bin  $B_i$
- $C(B_1) + C(B_m) > 1$  (b)(Otherwise, the items in  $B_m$  will be put in  $B_1$ .)
- For  $m$  nonempty bins,  
 $C(B_1) + C(B_2) + \dots + C(B_m) > m/2$ , (a)+(b) for  $i=1, \dots, m$   
 $\Rightarrow FF = m < 2 \sum_{i=1}^m C(B_i) = 2 \sum_{i=1}^n a_i \leq 2 OPT$   
 $FF < 2 OPT$

## Load balancing

**Input.**  $m$  identical machines;  $n \geq m$  jobs, job  $j$  has processing time  $t_j$ .


- Job  $j$  must run contiguously on one machine.
- A machine can process at most one job at a time.

**Def.** Let  $S[i]$  be the subset of jobs assigned to machine  $i$ .

The **load** of machine  $i$  is  $L[i] = \sum_{j \in S[i]} t_j$ .

**Def.** The **makespan** is the maximum load on any machine  $L = \max_i L[i]$ .

**Load balancing.** Assign each job to a machine to minimize makespan.

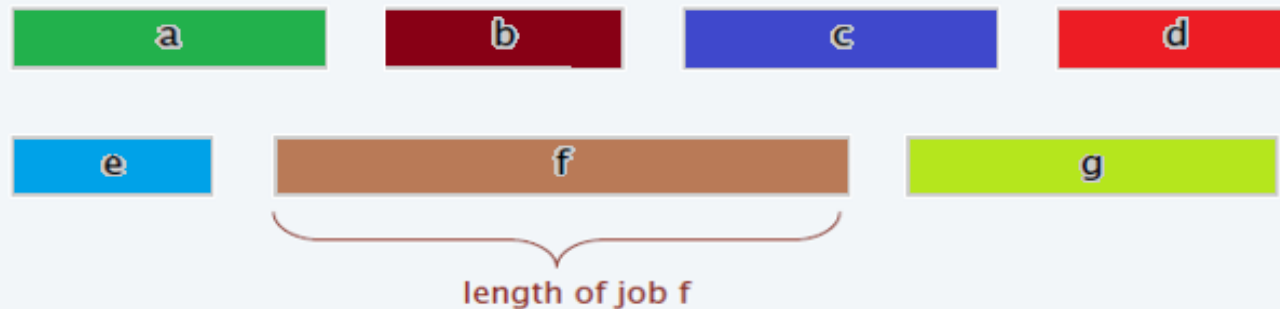


# Load balancing on 2 machines is NP-Hard/NP-Complete Problem

**Claim.** Load balancing is hard even if  $m = 2$  machines.

**Pf.**  $\text{PARTITION} \leq_p \text{LOAD-BALANCE}$ .

NP-Complete Optimization Problem



machine 1



machine 2



yes



## Load balancing: list scheduling

### List-scheduling algorithm.

- Consider  $n$  jobs in some fixed order.
- Assign job  $j$  to machine  $i$  whose load is smallest so far.

LIST-SCHEDULING ( $m, n, t_1, t_2, \dots, t_n$ )

---

FOR  $i = 1$  TO  $m$

$L[i] \leftarrow 0.$      $\leftarrow$  load on machine  $i$

$S[i] \leftarrow \emptyset.$      $\leftarrow$  jobs assigned to machine  $i$

FOR  $j = 1$  TO  $n$

$i \leftarrow \operatorname{argmin}_k L[k].$      $\leftarrow$  machine  $i$  has smallest load

$S[i] \leftarrow S[i] \cup \{j\}.$      $\leftarrow$  assign job  $j$  to machine  $i$

$L[i] \leftarrow L[i] + t_j.$      $\leftarrow$  update load of machine  $i$

RETURN  $S[1], S[2], \dots, S[m].$

Implementation.  $O(n \log m)$  using a priority queue for loads  $L[k]$ .

## Load balancing: list scheduling analysis

---

**Theorem.** [Graham 1966] Greedy algorithm is a 2-approximation.

- First worst-case analysis of an approximation algorithm.
- Need to compare resulting solution with optimal makespan  $L^*$ .

**Lemma 1.** For all  $k$ : the optimal makespan  $L^* \geq t_k$ .

**Pf.** Some machine must process the most time-consuming job. ■

**Lemma 2.** The optimal makespan  $L^* \geq \frac{1}{m} \sum_k t_k$ .

**Pf.**

- The total processing time is  $\sum_k t_k$ .
- One of  $m$  machines must do at least a  $1 / m$  fraction of total work. ■


## Load balancing: list scheduling analysis

**Theorem.** Greedy algorithm is a 2-approximation.

**Pf.** Consider load  $L[i]$  of bottleneck machine  $i$ . ← machine that ends up with highest load


- Let  $j$  be last job scheduled on machine  $i$ .
- When job  $j$  assigned to machine  $i$ ,  $i$  had smallest load.

Its load before assignment is  $L[i] - t_j$ ; hence  $L[i] - t_j \leq L[k]$  for all  $1 \leq k \leq m$ .



## Load balancing: list scheduling analysis

**Theorem.** Greedy algorithm is a 2-approximation.

**Pf.** Consider load  $L[i]$  of bottleneck machine  $i$ .  machine that ends up with highest load

- Let  $j$  be last job scheduled on machine  $i$ .
- When job  $j$  assigned to machine  $i$ ,  $i$  had smallest load.

Its load before assignment is  $L[i] - t_j$ ; hence  $L[i] - t_j \leq L[k]$  for all  $1 \leq k \leq m$ .

- Sum inequalities over all  $k$  and divide by  $m$ :

$$\begin{aligned} L[i] - t_j &\leq \frac{1}{m} \sum_k L[k] \\ &= \frac{1}{m} \sum_k t_k \end{aligned}$$

Lemma 2   $\leq L^*$ .

- Now,  $L = L[i] = (L[i] - t_j) + t_j \leq 2L^*$  .

$$\begin{array}{ccc} \underbrace{\phantom{L[i] - t_j}}_{\leq L^*} & + & \underbrace{\phantom{t_j}}_{\leq L^*} \\ \uparrow & & \uparrow \\ \text{above inequality} & & \text{Lemma 1} \end{array}$$



## Load balancing: list scheduling analysis

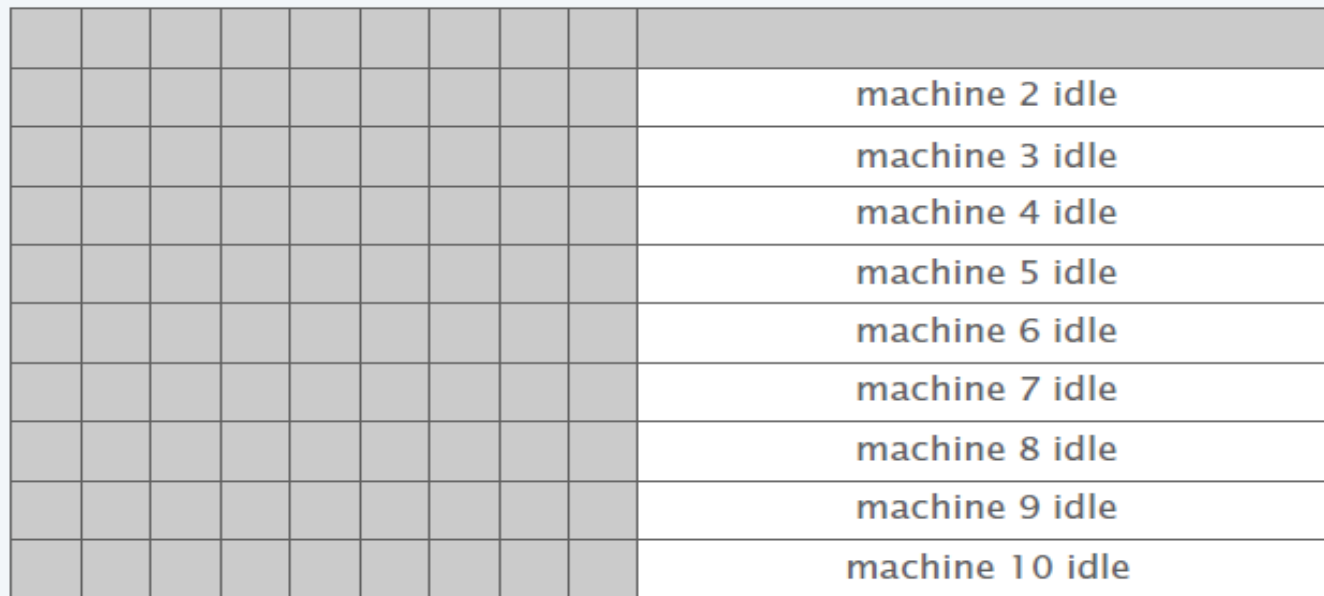
Q. Is our analysis tight?

A. Essentially yes.

Ex:  $m$  machines, first  $m(m-1)$  jobs have length 1, last job has length  $m$ .

list scheduling makespan =  $19 = 2m - 1$

$m = 10$



## Load balancing: list scheduling analysis


Q. Is our analysis tight?

A. Essentially yes.

Ex:  $m$  machines, first  $m(m-1)$  jobs have length 1, last job has length  $m$ .

optimal makespan =  $10 = m$

$m = 10$



## Load balancing: LPT rule

Longest processing time (LPT). Sort  $n$  jobs in decreasing order of processing times; then run list scheduling algorithm.

LPT-LIST-SCHEDULING ( $m, n, t_1, t_2, \dots, t_n$ )

SORT jobs and renumber so that  $t_1 \geq t_2 \geq \dots \geq t_n$ .

FOR  $i = 1$  TO  $m$

$L[i] \leftarrow 0.$   $\leftarrow$  load on machine  $i$

$S[i] \leftarrow \emptyset.$   $\leftarrow$  jobs assigned to machine  $i$

FOR  $j = 1$  TO  $n$

$i \leftarrow \operatorname{argmin}_k L[k].$   $\leftarrow$  machine  $i$  has smallest load

$S[i] \leftarrow S[i] \cup \{j\}.$   $\leftarrow$  assign job  $j$  to machine  $i$

$L[i] \leftarrow L[i] + t_j.$   $\leftarrow$  update load of machine  $i$

RETURN  $S[1], S[2], \dots, S[m].$

## Load balancing: LPT rule

**Observation.** If bottleneck machine  $i$  has only 1 job, then optimal.

**Pf.** Any solution must schedule that job. ■

**Lemma 3.** If there are more than  $m$  jobs,  $L^* \geq 2t_{m+1}$ .

**Pf.**

- Consider processing times of first  $m+1$  jobs  $t_1 \geq t_2 \geq \dots \geq t_{m+1}$ .
- Each takes at least  $t_{m+1}$  time.
- There are  $m+1$  jobs and  $m$  machines, so by pigeonhole principle, at least one machine gets two jobs. ■

**Theorem.** LPT rule is a  $3/2$ -approximation algorithm.

**Pf.** [ similar to proof for list scheduling ]

- Consider load  $L[i]$  of bottleneck machine  $i$ .
- Let  $j$  be last job scheduled on machine  $i$ . ← assuming machine  $i$  has at least 2 jobs, we have  $j \geq m+1$

$$L = L[i] = \underbrace{(L[i] - t_j)}_{\text{as before}} + \underbrace{t_j}_{\leq \frac{1}{2} L^*} \leq \frac{3}{2} L^* \quad \blacksquare$$

as before  $\rightarrow \leq L^* \quad \leq \frac{1}{2} L^* \leftarrow \text{Lemma 3 (since } t_{m+1} \geq t_j)$

## Load balancing: LPT rule

---

Q. Is our  $3/2$  analysis tight?

A. No.

Theorem. [Graham 1969] LPT rule is a  $4/3$ -approximation.

Pf. More sophisticated analysis of same algorithm.

Q. Is Graham's  $4/3$  analysis tight?

A. Essentially yes.

Ex.

- $m$  machines
- $n = 2m + 1$  jobs
- 2 jobs of length  $m, m + 1, \dots, 2m - 1$  and one more job of length  $m$ .
- Then,  $L / L^* = (4m - 1) / (3m)$

## Generalized load balancing

---

**Input.** Set of  $m$  machines  $M$ ; set of  $n$  jobs  $J$ .

- Job  $j \in J$  must run contiguously on an **authorized machine** in  $M_j \subseteq M$ .
- Job  $j \in J$  has processing time  $t_j$ .
- Each machine can process at most one job at a time.

**Def.** Let  $J_i$  be the subset of jobs assigned to machine  $i$ .

The load of machine  $i$  is  $L_i = \sum_{j \in J_i} t_j$ .

**Def.** The makespan is the maximum load on any machine  $= \max_i L_i$ .

**Generalized load balancing.** Assign each job to an authorized machine to minimize makespan.

## Generalized load balancing: integer linear program and relaxation

ILP formulation.  $x_{ij}$  = time machine  $i$  spends processing job  $j$ .

$$\begin{aligned} (IP) \quad & \min \quad L \\ & \text{s. t.} \quad \sum_i x_{ij} = t_j \quad \text{for all } j \in J \\ & \quad \sum_j x_{ij} \leq L \quad \text{for all } i \in M \\ & \quad x_{ij} \in \{0, t_j\} \quad \text{for all } j \in J \text{ and } i \in M_j \\ & \quad x_{ij} = 0 \quad \text{for all } j \in J \text{ and } i \notin M_j \end{aligned}$$

LP relaxation.

$$\begin{aligned} (LP) \quad & \min \quad L \\ & \text{s. t.} \quad \sum_i x_{ij} = t_j \quad \text{for all } j \in J \\ & \quad \sum_j x_{ij} \leq L \quad \text{for all } i \in M \\ & \quad x_{ij} \geq 0 \quad \text{for all } j \in J \text{ and } i \in M_j \\ & \quad x_{ij} = 0 \quad \text{for all } j \in J \text{ and } i \notin M_j \end{aligned}$$

# Generalized load balancing: lower bounds

---

Lemma 1. The optimal makespan  $L^* \geq \max_j t_j$ .

Pf. Some machine must process the most time-consuming job. ■

Lemma 2. Let  $L$  be optimal value to the  $LP$ . Then, optimal makespan  $L^* \geq L$ .

Pf.  $LP$  has fewer constraints than  $ILP$  formulation. ■




## Generalized load balancing: structure of LP solution

**Lemma 3.** Let  $x$  be solution to  $LP$ . Let  $G(x)$  be the graph with an edge between machine  $i$  and job  $j$  if  $x_{ij} > 0$ . Then  $G(x)$  is **acyclic**.


**Pf.** (deferred)

can transform  $x$  into another LP solution where  $G(x)$  is acyclic if LP solver doesn't return such an  $x$



$G(x)$  acyclic

$x_{ij} > 0$



$G(x)$  cyclic

○ job

□ machine

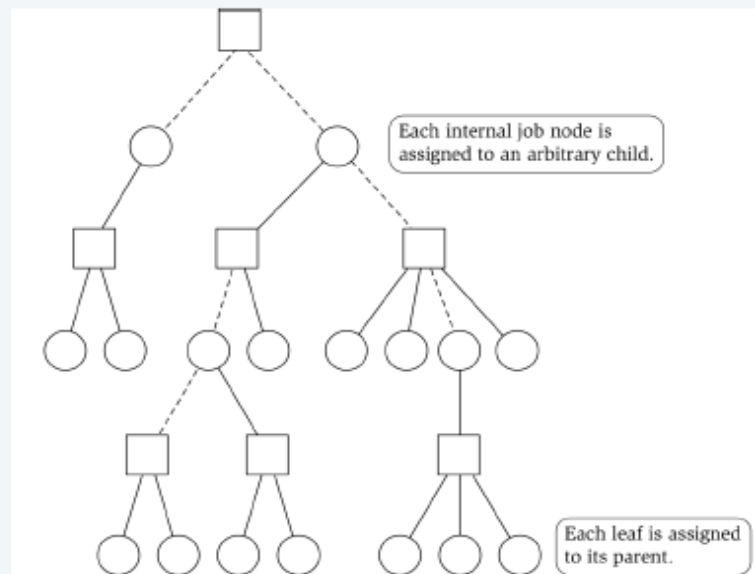

## Generalized load balancing: rounding

**Rounded solution.** Find  $LP$  solution  $x$  where  $G(x)$  is a forest. Root forest  $G(x)$  at some arbitrary machine node  $r$ .

- If job  $j$  is a leaf node, assign  $j$  to its parent machine  $i$ .
- If job  $j$  is not a leaf node, assign  $j$  to any one of its children.

**Lemma 4.** Rounded solution only assigns jobs to authorized machines.

**Pf.** If job  $j$  is assigned to machine  $i$ , then  $x_{ij} > 0$ .  $LP$  solution can only assign positive value to authorized machines. ■



## Generalized load balancing: analysis

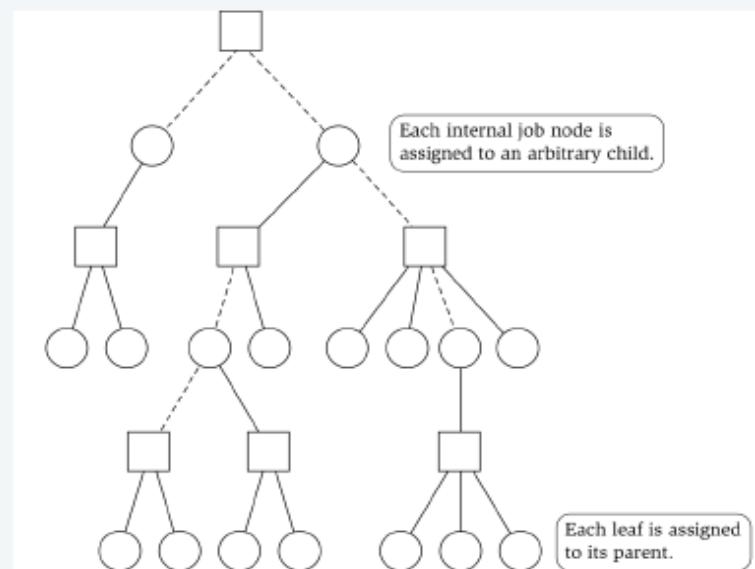

**Lemma 5.** If job  $j$  is a leaf node and machine  $i = \text{parent}(j)$ , then  $x_{ij} = t_j$ .

**Pf.**

- Since  $i$  is a leaf,  $x_{ij} = 0$  for all  $j \neq \text{parent}(i)$ .
- LP constraint guarantees  $\sum_i x_{ij} = t_j$ . ■

**Lemma 6.** At most one non-leaf job is assigned to a machine.

**Pf.** The only possible non-leaf job assigned to machine  $i$  is  $\text{parent}(i)$ . ■



## Generalized load balancing: analysis

**Theorem.** Rounded solution is a 2-approximation.

**Pf.**

- Let  $J(i)$  be the jobs assigned to machine  $i$ .
- By LEMMA 6, the load  $L_i$  on machine  $i$  has two components:

- leaf nodes:

$$\sum_{\substack{j \in J(i) \\ j \text{ is a leaf}}} t_j \stackrel{\text{Lemma 5}}{=} \sum_{\substack{j \in J(i) \\ j \text{ is a leaf}}} x_{ij} \leq \sum_{j \in J} x_{ij} \leq L \stackrel{\text{LP}}{\leq} L^* \stackrel{\text{Lemma 2 (LP is a relaxation)}}{\leq} L^*$$

optimal value of LP

- parent:  $t_{\text{parent}(i)} \stackrel{\text{Lemma 1}}{\leq} L^*$

- Thus, the overall load  $L_i \leq 2L^*$ . ■

# Generalized load balancing: flow formulation


## Flow formulation of $LP$ .

$$\sum_i x_{ij} = t_j \quad \text{for all } j \in J$$

$$\sum_j x_{ij} \leq L \quad \text{for all } i \in M$$

$$x_{ij} \geq 0 \quad \text{for all } j \in J \text{ and } i \in M_j$$

$$x_{ij} = 0 \quad \text{for all } j \in J \text{ and } i \notin M_j$$





**Observation.** Solution to feasible flow problem with value  $L$  are in 1-to-1 correspondence with  $LP$  solutions of value  $L$ .

## Generalized load balancing: structure of solution

**Lemma 3.** Let  $(x, L)$  be solution to  $LP$ . Let  $G(x)$  be the graph with an edge from machine  $i$  to job  $j$  if  $x_{ij} > 0$ . We can find another solution  $(x', L)$  such that  $G(x')$  is acyclic.

**Pf.** Let  $C$  be a cycle in  $G(x)$ .

- Augment flow along the cycle  $C$ .  $\leftarrow$  flow conservation maintained
- At least one edge from  $C$  is removed (and none are added).
- Repeat until  $G(x')$  is acyclic. ■



## Conclusions

---

**Running time.** The bottleneck operation in our 2-approximation is solving one  $LP$  with  $mn + 1$  variables.

**Remark.** Can solve  $LP$  using flow techniques on a graph with  $m + n + 1$  nodes: given  $L$ , find feasible flow if it exists. Binary search to find  $L^*$ .

**Extensions: unrelated parallel machines.** [Lenstra–Shmoys–Tardos 1990]

- Job  $j$  takes  $t_{ij}$  time if processed on machine  $i$ .
- 2-approximation algorithm via LP rounding.
- If  $P \neq NP$ , then no  $\rho$ -approximation exists for any  $\rho < 3/2$ .